

Honors Machine Learning

# **DYNAMIC PROGRAMMING**

## in Reinforcement Learning

*Jaxon Ham & Ruth Walters*

# Reward and Return

- **reward:** the feedback that the agent receives from interaction with the environment
  - $R_{t+1}$  is the immediate reward
  - $R_{t+2}, R_{t+3}, \dots R_{t+n}$  are the subsequent  $n$  rewards
- **return:** the cumulative reward obtained over the duration of an episode
- **discount factor:** ( $\gamma$ ) the extent to which future rewards are valued

## Reward and Return

The return  $G_t$  at time  $t$  is calculated from the immediate return and the weighted sum of all subsequent returns:

**Weighted sum:**  $G_t \stackrel{\text{def}}{=} R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots$

$$= \sum_{k=0} \gamma^k R_{t+k+1}$$

**Recursive:** 
$$\begin{aligned} G_t &= R_{t+1} + \gamma G_{t+1} \\ &= r + \gamma G_{t+1} \end{aligned}$$

# Policy

- **policy:**  $\pi(a | s)$  a function that determines the next action  $a$  to take given a specific state  $s$
- **optimal policy:**  $\pi_*(a | s)$  the policy that yields the highest return
- In reinforcement learning, agents start with a random policy that they update to maximize return as they gain experience interacting with their environment

# Value Function

- **value function:** function that measures the favorability of each state based on the return  $G_t$

The value function  $v_\pi(s)$  of state  $s$  represents the expected return of a policy,  $\pi$ :

$$\begin{aligned} v_\pi(s) &\stackrel{\text{def}}{=} E_\pi[G_t \mid S_t = s] \\ &= E_\pi \left[ \sum_{k=0} \gamma^{k+1} R_{t+k+1} \mid S_t = s \right] \end{aligned}$$

## Action Value Function

- **action value function:** the expected return  $G_t$  of action  $A_t = a$  when the agent is in state  $S_t = s$

The action value function  $q_\pi(s, a)$  of state  $s$  and action  $a$  represents the expected return of a policy,  $\pi$  for  $a$  and  $s$


$$\begin{aligned} q_\pi(s, a) &\stackrel{\text{def}}{=} E_\pi[G_t \mid S_t = s, A_t = a] \\ &= E_\pi \left[ \sum_{k=0} \gamma^{k+1} R_{t+k+1} \mid S_t = s, A_t = a \right] \end{aligned}$$

# Summary


1. The **policy** decides an action
2. Each action produces a **reward**, which prompts the transition from one state to another. A reward is a *consequence* of an action.
3. The **return** is the stream of increasingly discounted future rewards. The return is the *cumulation* of rewards.
4. The **value function** is the expected return from a state under the policy
5. The **action-value function** is the expected return from a state after choosing a specific action and following the policy

# Reinforcement Learning Algorithms


When the **dynamics** of the environment are **known**, a learning task can be solved with dynamic programming



# Model-Free methods



# Model-Based methods, e.g. dynamic programming



Dynamic programming models don't have to interact with their environment

## *the Bellman Equation*

The value function for a state  $s$  can be given as a function of its subsequent state,  $s'$ :

$$v_{\pi}(s) = \sum_{a \in \hat{A}} \pi(a \mid s) \sum_{s' \in \hat{S}, r \in \hat{R}} p(s', r \mid s, a) [r + \gamma v_{\pi}(s')]$$

# **Assumptions *in* Dynamic Programming**

1

The environment dynamics are known

2

The agent's state has the Markov property

- **Markov property:** a property of an agent's state in which **the next action and reward depend *only* on the current state and choice of action** made at the current moment (time step)

# Objectives of a Dynamic Programming Model

1

## Prediction Task

Obtain the true-state value function,  $v_\pi(s)$ , of a policy  $\pi$  via **policy evaluation**

2

## Policy Improvement Task

Find the optimal value function,  $v_*(s)$ , of a policy  $\pi$  via **generalized policy iteration**

Value Iteration

# Policy Evaluation

1. Start at  $v^{(0)}(s)$ , with zero values for each state
2. At each iteration, update the values for each state using the Bellman equation

$$v^{(i+1)}(s) = \sum_a \pi(a \mid s) \sum_{s' \in \hat{S}, r \in \hat{R}} p(s', r \mid s, a)[r + \gamma v^{(i)}(s')]$$

- As  $i$  increases towards infinity, the value function  $v^{(i+1)}(s)$  will converge to the true state value function
- Since the dynamics of the environment are known, there is no need to interact with the environment

# Policy Iteration

- Once the value function  $v_\pi(s)$  has been determined for the current policy  $\pi$ , the value function can be used to improve that policy
- The **goal** is to find a new policy  $\pi'$  that, **for each state  $s$**  would yield **higher or at least equal value** relative to the current policy

$$v_{\pi'}(s) \geq v_\pi(s) \quad \forall s \in \hat{S}$$

# Generalized Policy Improvement

## Iterate

1. Compute the action-value function  $q_{\pi}(s, a)$  for each state  $s$  and each action  $a$
2. For each  $q_{\pi}(s, a)$ , compare the value if the next state  $s'$  if action  $a$  was selected

## Evaluate

Compare the action  $a_*$  that produces the highest state value to the action selected by the current policy

## Update

If the action suggested by the current policy is worse than the action suggested by the action value-function, update the policy

# Value Iteration

- **value iteration:** the process of selecting the best action out of all possible actions by combining policy evaluation and policy improvement steps
- Combines policy evaluation and policy iteration
- The updated value for  $v^{(i+1)}(s)$  is maximized by choosing the best action out of all possible actions

# Value Iteration

- The value function  $v^{(i+1)}$  for iteration  $i + 1$  is updated based on the action that maximizes the weighted sum of the next state value and its immediate reward

$$v^{(i+1)}(s) = \max_a \sum_{s', r} p(s', r | s, a)[r + \gamma v^{(i)}(s')]$$

- Compare to policy evaluation, which updates using a weighted sum of all actions


# Example Context



**Basic Concept: Warehouse Robots**

# Building the “GridWorld” Environment

- **Rewards (R):**
  - -1 per move (step cost)
  - -2 to -6 for hazards
  - +10 for reaching goal
- **Transition Model (P):** 10% slip chance → imperfect movement
- **Discount ( $\gamma = 0.97$ ):** favors quicker paths



# Implementing the Bellman Equation

**Bellman Optimality Equation:** 
$$V(s) = \max_a \sum_{s',r} P(s',r|s,a)[r + \gamma V(s')]$$

**In Code from `value_iteration()`:**


- $V(s)$ : state's value
- $a$ : action
- $s'$ : next state
- $r$ : immediate reward
- $P(s', r | s, a)$ : transition probability

```
q = [sum(p * (r + gamma * V[s2])
         for p, s2, r, d in P[s][a]) for a in ACTIONS]
V[s] = max(q)
```

State  $s \rightarrow$  Action  $a \rightarrow$  Next state  $s' \rightarrow$  Reward  $r \rightarrow$  Updated Value  $V(s)$


# The Dynamic Programming Process

- **Initialize:** all  $V(s) = 0$
- **Update:** apply Bellman equation repeatedly
- **Improve:** pick best action per state (greedy)
- **Converge:** stop when  $\Delta V <$  threshold
- Combines:
  - **Policy Evaluation** (find  $V\pi$ )
  - **Policy Improvement** (make  $\pi$  greedy)



# Value Heatmap: Learned State Values

- Each color = expected total reward  $V^*(s)$
- Brighter → higher value (closer to goal)
- Darker → lower value (hazards / farther away)
- DP propagates value backward from the goal



# Policy Grid: Extract Optimal Actions

**Policy Formula:**


$$\pi^*(s) = \arg \max_a \sum_{s', r} P(s', r | s, a) [r + \gamma V^*(s')]$$

**In Code**

**greedy\_policy\_from\_V():**

- Arrows = best actions per state
- Hazards cause arrows to reroute

```
pi[s, np.argmax(q)] = 1
```



# Rollout Simulation: Testing the Policy

- Simulates following  $\pi^*$  from Start → Goal
- Each step:
  - Pick best action:  $\text{argmax}(\pi[s])$
  - Random 10% slip chance
  - Add step/hazard/goal reward
- **Rollout path:** sequence of states
- **Total return:** sum of all rewards

